


Homework 1 Report - PM2.5 Prediction

學號：R07922135 系級：資工研一 姓名：顏百謙

助教您好：我是有和您說我們要出國參加研討會的同學，所以 1-3 題的內容都只有用 public 來表示。

1. (1%) 請分別使用至少 4 種不同數值的 **learning rate** 進行 **training** (其他參數需一致)，對其作圖，並且討論其收斂過程差異。



在 train 每小時氣溫，臭氧，pm10，pm2.5 的狀況下以

1. 綠線 learning rate = 0.0000001
2. 藍線 learning rate = 0.000001
3. 紅線 learning rate = 0.000008
4. 黑線 learning rate = 0.0000208

Train 200 個 iteration，並對其 RMSE 變化作圖，由圖可知當 learning rate 適量增加時可有效提昇收斂速度，但若 learning rate 過高會反而造成結果變差。

2. (1%) 請分別使用每筆 **data9** 小時內所有 **feature** 的一次項 (含 **bias** 項) 以及每筆 **data9** 小時內 **PM2.5** 的一次項 (含 **bias** 項) 進行 **training**, 比較並討論這兩種模型的 **root mean-square error** (根據 **kaggle** 上的 **public/private score**) 。

使用處理過的training data / learning rate = 0.000001 / iteration = 20000 進行training 得到

(1.) 僅使用pm2.5的model得到了8.76149 Score

Name	Submitted	Wait time	Execution time	Score
hw1-feature-pm25.csv	just now	0 seconds	0 seconds	8.76149
Complete				
Jump to your position on the leaderboard ▾				

(2.) 使用了全部feature 的model 得到了23.24591 Score

Name	Submitted	Wait time	Execution time	Score
hw1-feature-all.csv	just now	0 seconds	0 seconds	23.24591
Complete				
Jump to your position on the leaderboard ▾				

3. (1%)請分別使用至少四種不同數值的 **regularization parameter** λ 進行 **training** (其他參數需一至) , 討論及討論其 **RMSE(traning, testing)** (testing 根據 kaggle 上的 **public/private score**) 以及參數 **weight** 的 **L2 norm**。

使用處理過的training data / learning rate = 0.000001 / iteration = 20000 / feature = 9 小時內每小時的「氣溫」, 「臭氧」, 「pm10」, 「pm2.5」進行training 得到

$\lambda = 0$ 時得到 9.08804 Score 且weight 的 L2 norm = 0.9350074747035682

Name	Submitted	Wait time	Execution time	Score
hw1-randa-0.csv	just now	0 seconds	0 seconds	9.08804
Complete				
Jump to your position on the leaderboard ▾				

$\lambda = 1000$ 時得到 9.04907 Score 且weight 的 L2 norm = 0.9326304150157435

Name	Submitted	Wait time	Execution time	Score
hw1-randa-1000.csv	just now	0 seconds	0 seconds	9.04907
Complete				
Jump to your position on the leaderboard ▾				

$\lambda = 100000$ 時得到 8.77764 Score 且weight 的 L2 norm = 0.7476988271218137

Name	Submitted	Wait time	Execution time	Score
hw1-randa-100000.csv	just now	0 seconds	0 seconds	8.77764
Complete				
Jump to your position on the leaderboard ▾				

$\lambda = 1000000$ 時得到9.68059 Score 且weight 的 L2 norm = 0.3520363257057711

Name	Submitted	Wait time	Execution time	Score
hw1-randa-1000000.csv	just now	0 seconds	0 seconds	9.68059
Complete				
Jump to your position on the leaderboard ▾				

由結果可知, 在適度增加 λ 時可以使weight 趨向小值, 且使結果變好, 但當 λ 過大時會造成 x 的影響力過小(幾乎train 不到) 反而使 RMSE 結果變差

4 (1%)

(4-a)

Given t_n is the data point of the data set $\mathcal{D} = \{t_1, \dots, t_N\}$. Each data point t_n is associated with a weighting factor $r_n > 0$. The sum-of-squares error function becomes:

$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N r_n (t_n - \mathbf{w}^T \mathbf{x}_n)^2$ Find the solution \mathbf{w}^* that minimizes the error function.

$$\text{Let } (\sqrt{r_1}t_1, \sqrt{r_2}t_2, \dots) = \hat{t}, \quad \begin{bmatrix} x_1^1 \sqrt{r_1} & x_2^1 \sqrt{r_2} & \dots & x_n^1 \sqrt{r_n} \\ x_1^2 \sqrt{r_1} & x_2^2 \sqrt{r_2} & \dots & x_n^2 \sqrt{r_n} \\ \vdots & \vdots & \ddots & \vdots \\ x_1^m \sqrt{r_1} & x_2^m \sqrt{r_2} & \dots & x_n^m \sqrt{r_n} \end{bmatrix} = \hat{x}$$

$$\begin{aligned} E_D(w) &= \frac{1}{2} \sum_{n=1}^N (\sqrt{r_n}t_n - \sqrt{r_n}w^T x_n)^2 \\ &= \frac{1}{2} (\hat{t} - w^T \hat{x})(\hat{t} - w^T \hat{x})^T \\ &= \frac{1}{2} (\hat{t} \hat{t}^T - w^T \hat{x} \hat{t}^T - \hat{t} \hat{x}^T w + w^T \hat{x} \hat{x}^T w) \end{aligned}$$

Find w^* to minimize $E_D(w) = \text{Find } \nabla E_D(w^*) = 0$

$$\nabla E_D(w^*) = -\hat{x} \hat{t}^T + \hat{x} \hat{x}^T w$$

So, when $w^* = (\hat{x} \hat{x}^T)^{-1} (\hat{x} \hat{t}^T)$ can minimize $E_D(w)$

(4-b)

Following the previous problem(2-a), if

$\mathbf{t} = [t_1 t_2 t_3] = [0 \quad 10 \quad 5]$, $\mathbf{X} = [\mathbf{x}_1 \mathbf{x}_2 \mathbf{x}_3] = \begin{bmatrix} 2 & 5 & 5 \\ 3 & 1 & 6 \end{bmatrix}$ $r_1 = 2, r_2 = 1, r_3 = 3$ Find the solution \mathbf{w}^* .

$$\hat{x} = \begin{bmatrix} 2\sqrt{2} & 5 & 5\sqrt{3} \\ 3\sqrt{2} & 1 & 6\sqrt{3} \end{bmatrix}, \quad \hat{t} = [0 \quad 10 \quad 5\sqrt{3}]$$

$$\begin{aligned} w^* &= \left(\begin{bmatrix} 2\sqrt{2} & 5 & 5\sqrt{3} \\ 3\sqrt{2} & 1 & 6\sqrt{3} \end{bmatrix} \begin{bmatrix} 2\sqrt{2} & 3\sqrt{2} \\ 5 & 1 \\ 5\sqrt{3} & 6\sqrt{3} \end{bmatrix} \right)^{-1} \left(\begin{bmatrix} 125 \\ 100 \end{bmatrix} \right) \\ &= \begin{bmatrix} 108 & 107 \\ 107 & 127 \end{bmatrix}^{-1} \begin{bmatrix} 125 \\ 100 \end{bmatrix} \\ &= \begin{bmatrix} \frac{5175}{2267} \\ -\frac{2575}{2267} \end{bmatrix} \end{aligned}$$

5 (1%)

Given a linear model:

$$y(x, \mathbf{w}) = w_0 + \sum_{i=1}^D w_i x_i$$

with a sum-of-squares error function:

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (y(x_n, \mathbf{w}) - t_n)^2$$

where t_n is the data point of the data set $\mathcal{D} = \{t_1, \dots, t_N\}$

Suppose that Gaussian noise ϵ_i with zero mean and variance σ^2 is added independently to each of the input variables x_i . By making use of $\mathbb{E}[\epsilon_i \epsilon_j] = \delta_{ij} \sigma^2$ and $\mathbb{E}[\epsilon_i] = 0$, show that minimizing E averaged over the noise distribution is equivalent to minimizing the sum-of-squares error for noise-free input variables with the addition of a weight -decay regularization term, in which the bias parameter w_0 is omitted from the regularizer.

Hint

- $\delta_{ij} = \begin{cases} 1(i = j), \\ 0(i \neq j). \end{cases}$

$$\begin{aligned} \tilde{y}(x, w) &= w_0 + \sum_{i=1}^D (w_i x_i) \\ \tilde{E}(w) &= \frac{1}{2} \sum_{n=1}^N (w_0 + \sum_{i=1}^D (w_i x_{i,n}) - t_n)^2 + \lambda \sum_{i=1}^D w_i^2 \\ &= \frac{1}{2} \sum_{n=1}^N (\tilde{y}_n^2 - 2\tilde{y}_n t_n + t_n^2) + \lambda \sum_{i=1}^D w_i^2 \end{aligned}$$

$$\hat{y}(x, w) = w_0 + \sum_{i=1}^D (w_i x_i + \epsilon_i)$$

$$\begin{aligned}\hat{E}(w) &= \frac{1}{2} \sum_{n=1}^N (w_0 + \sum_{i=1}^D (w_i x_{i,n} + \epsilon_{i,n}) - t_n)^2 \\ &= \frac{1}{2} \sum_{n=1}^N (\tilde{y}_n^2 - 2\tilde{y}_n t_n + t_n^2) \\ &= \frac{1}{2} \sum_{n=1}^N (\tilde{y}_n^2 + (\epsilon_{1,n} + \epsilon_{2,n} + \dots)(w_0 + w_1 x_{1,n} + \dots) + (\epsilon_{1,n}^2 + \epsilon_{2,n}^2 + \dots) \\ &\quad + (\sum_{i=1}^D \sum_{j=i+1}^D \epsilon_{i,n} \epsilon_{j,n}) + 2\tilde{y}_n t_n + (t_n \sum_{i=1}^D \epsilon_{i,n}) + t_n^2)\end{aligned}$$

$$\text{because } \sum_{i=1}^D \sum_{j=i+1}^D \epsilon_{i,n} \epsilon_{j,n} = 0, \sum_{i=1}^D \epsilon_i = 0 \quad \text{we can get}$$

$$\begin{aligned}&= \frac{1}{2} \sum_{n=1}^N (\tilde{y}_n^2 + (\epsilon_{1,n}^2 + \epsilon_{2,n}^2 + \dots) + 2\tilde{y}_n t_n + t_n^2) \\ &= \frac{1}{2} \sum_{n=1}^N (\tilde{y}_n^2 - 2\tilde{y}_n t_n + t_n^2) + \frac{1}{2} \sum_{n=1}^N (\epsilon_{1,n}^2 + \epsilon_{2,n}^2 + \dots) \\ &= \frac{1}{2} \sum_{n=1}^N (\tilde{y}_n^2 - 2\tilde{y}_n t_n + t_n^2) + \frac{N}{2} \sigma^2\end{aligned}$$

因此當我們將 λ 令為 $\frac{N\sigma^2}{2 \sum_{i=1}^D w_i^2}$ 可獲得相同的效果

6 (1%)

$\mathbf{A} \in \mathbb{R}^{n \times n}$, α is one of the elements of \mathbf{A} , prove that

$\frac{d}{d\alpha} \ln|\mathbf{A}| = \text{Tr}\left(\mathbf{A}^{-1} \frac{d}{d\alpha} \mathbf{A}\right)$ where the matrix \mathbf{A} is a real, symmetric, non-singular matrix.

Hint:

- The determinant and trace of \mathbf{A} could be expressed in terms of its eigenvalues.

A is symmetric matrix

$$A \text{ 可分解為 } PDP^{-1}$$

$$\text{其中 } D = \begin{bmatrix} \lambda_1 & 0 & \cdots & 0 \\ 0 & \lambda_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \lambda_n \end{bmatrix}$$

$$\det(e^D) = e^{\lambda_1} e^{\lambda_2} e^{\lambda_3} \cdots e^{\lambda_n} = e^{\text{Tr}(D)}$$

$$\text{Because } A^k = PD^k P^{-1}, \text{ for all } k$$

$$\text{we can get } A^{\ln(e)} = PD^{\ln(e)} P^{-1}$$

$$e^{\ln(A)} = P e^{\ln(D)} P^{-1}$$

$$\text{and } \text{Tr}(A) = \text{Tr}(PDP^{-1}) = \text{Tr}(D)$$

$$\text{so } \det(e^{\ln(A)}) = \det(e^{\ln(D)}) = e^{\text{Tr}(\ln(D))} = e^{\text{Tr}(\ln(A))}$$

$$\det(A) = e^{\text{Tr}(\ln(A))}$$

$$\ln(\det(A)) = \text{Tr}(\ln(A))$$

$$\begin{aligned} \frac{d}{d\alpha} \ln|A| &= \frac{d}{d\alpha} \text{Tr}(\ln(A)) \\ &= \text{Tr}\left(\frac{d}{d\alpha} \ln(A)\right) \quad , \text{By } \left(\frac{\partial g(u)}{\partial x} = \frac{\partial g(u)}{\partial u} \frac{\partial u}{\partial x}\right) \\ &= \text{Tr}\left(\frac{\partial \ln(A)}{\partial A} \frac{\partial A}{\partial \alpha}\right) \\ &= \text{Tr}\left(\frac{1}{A} \frac{dA}{d\alpha}\right) \end{aligned}$$